



A Review of Large-vocabulary Continuous-speech Recognition

Considerable progress has been made in speech-recognition technology over the last few years and nowhere has this progress been more evident than in the area of large-vocabulary recognition (LVR). Current laboratory systems are capable of transcribing continuous speech from any speaker with average word-error rates between 5% and 10%. If speaker adaptation is allowed, then after 2 or 3 minutes of speech, the error rate will drop well below 5% for most speakers.

LVR systems had been limited to dictation applications since the systems were speaker

dependent and required words to be spoken with a short pause between them. However, the capability to recognize natural continuous-speech input from any speaker opens up many more applications. As a result, LVR technology appears to be on the brink of widespread deployment across a range of information technology (IT) systems.

This article discusses the principles and architecture of current LVR systems and identifies the key issues affecting their future deployment. To illustrate the various points raised, the Cambridge University HTK system is described. This sys-

STEVE YOUNG

tem is a modern design that gives state-of-the-art performance, and it is typical of the current generation of recognition systems.

System Overview

Current LVR systems are firmly based on the principles of statistical pattern recognition. The basic methods of applying these principles to the problem of speech recognition were pioneered by Baker, Jelinek, and their colleagues from IBM in the 1970s, and little has changed since [13, 54]. Figure 1 illustrates the main components of an LVR system.

An unknown speech waveform is converted by a front-end signal processor into a sequence of acoustic vectors, $Y = y_1, y_2, \dots, y_T$. Each of these vectors is a compact representation of the short-time speech spectrum covering a period of typically 10 msecs. Thus, a typical 10-word utterance might have a duration of around 3 seconds and would be represented by a sequence of $T = 300$ acoustic vectors.

The utterance consists of a sequence of words, $W = w_1, w_2, \dots, w_n$, and it is the job of the LVR system to determine the most probable word sequence, \hat{W} , given the observed acoustic signal Y .


To do this, Bayes' rule is used to decompose the required probability $P(W|Y)$ into two components, that is,

$$\hat{W} = \arg \max_w P(W|Y) = \arg \max_w \frac{P(W)P(Y|W)}{P(Y)} \quad (1)$$


This equation indicates that to find the most likely word sequence W , the word sequence that maximizes the product of $P(W)$ and $P(Y|W)$ must be found. The first of these terms represents the *a priori* probability of observing W independent of the observed signal, and this probability is determined by a *language model*. The second term represents the probability of observing the vector sequence Y given some specified word sequence W , and this probability is determined by an *acoustic model*.

Figure 1 shows how these relationships might be computed. The word sequence $W = \text{"This is speech"}$ is postulated and the language model computes its probability $P(W)$. Each word is then converted into a sequence of basic sounds or *phones* using a pronouncing dictionary. For each phone there is a corresponding statistical model called a hidden Markov model (HMM). The sequence of HMMs needed to represent the postulated utterance are concatenated to form a single composite model, and the probability of that model generating the observed sequence Y is calculated. This is the required probability $P(Y|W)$. In principle, this process can be repeated for all possible word sequences with the most likely sequence selected as the recognizer output.

To convert the above design philosophy into a practical system requires the solution of a number of challenging problems. First, a front-end parameterization is needed that can extract from the speech waveform all of the necessary acoustic information in a compact form compatible with the HMM-based acoustic models. Second, the HMM models



1. Overview of Statistical Speech Recognition. This diagram shows the computation of the probability $P(W|Y)$ of word sequence W given the parameterized acoustic signal Y . The prior probability $P(W)$ is determined directly from a language model. The likelihood of the acoustic data $P(Y|W)$ is computed using a composite hidden Markov model representing W constructed from simple HMM phone models joined in sequence according to word pronunciations stored in a dictionary.



2. MFCC-based Front-End Processor. To perform pattern-matching, the speech waveform must be converted to a sequence of acoustic vectors representing a smoothed log spectrum computed every 10 msecs. Performance is improved by using a non-linear mel-frequency scale followed by a discrete cosine transform (DCT). The latter has the effect of decorrelating the signal, thereby improving assumptions of statistical independence. Finally, first and second differentials are appended to incorporate dynamic information about the signal.

themselves must accurately represent the distributions of each sound in each of the many contexts in which the sound may occur. Furthermore, the HMM parameters must be estimated from data, and it will never be possible to obtain sufficient data to cover all possible contexts. Third, the language model must be designed to give accurate word predictions based on the preceding history. However, as for the HMMs, data sparsity is an ever-present problem and the language model must be able to deal with word sequences for which no examples occur in the training data. Finally, the process outlined above for finding \hat{W} by enumerating all possible word sequences is clearly impractical. Instead, potential word sequences are explored in parallel, discarding hypotheses as soon as they become improbable. This process is called *decoding*, and the design of efficient decoders is crucial to the realization of practical LVR systems capable of fast and accurate operation on today's computing platforms. The next four sections of this article deal with each of these issues in more detail.


Front-End Parameterization

A key assumption made in the design of current recognizers is that the speech signal can be regarded as stationary (i.e., the spectral characteristics are relatively constant) over an interval of a few milliseconds. Thus, the prime function of the front-end parameterization stage is to divide the input speech into blocks and from each block derive a smoothed spectral estimate. The spacing between blocks is typically 10 msec, and blocks are normally overlapped to give a longer analysis window of typically 25 msec. As with all processing of this type, it is common to apply a tapered window function (e.g., Hamming) to each block. Also the speech signal is often pre-emphasized by applying high-frequency amplification to compensate for the attenuation caused by the radiation from the lips.

The required spectral estimates may be computed via linear prediction or Fourier analysis [89], and there are a number of additional transformations that can be applied in order to generate the final acoustic vectors. To illustrate one typical arrangement, Fig. 2 shows the front end of the HTK recognizer, which generates mel-frequency cepstral coefficients (MFCCs) [24].

To compute MFCC coefficients, the Fourier spectrum is smoothed by integrating the spectral coefficients within triangular frequency bins arranged on a non-linear scale called the *mel-scale*. For 8 kHz bandwidth speech, the HTK recognizer uses 24 of these triangular frequency bins. The mel-scale is designed to approximate the frequency resolution of the human ear and is linear up to 1000 Hz and logarithmic thereafter. More importantly, its use has been shown empirically to improve recognition accuracy [91]. In order to make the statistics of the estimated speech power spectrum approximately Gaussian, log compression is applied to the filter-bank output.

The final processing stage is to apply the discrete cosine transform (DCT) to the log filter-bank coefficients. This has



3. HMM-based Phone Model. An HMM can be regarded as a random generator of acoustic vectors. It consists of a sequence of states connected by probabilistic transitions. It changes to a new state each time period, generating a new acoustic vector according to the output distribution of that state. The transition probabilities therefore model the durational variability in real speech and the output probabilities model the spectral variability.

the effect of compressing the spectral information into the lower-order coefficients, and it also decorrelates them to allow the subsequent statistical modelling to use diagonal covariance matrices. In the HTK recognizer, the signal energy plus the first 12 cepstral coefficients are retained to form a basic 13-element acoustic vector. Cepstral coefficients can also be derived from LP coefficients where they achieve a similar decorrelating effect [4, 30]. Good results have also been reported using LP coefficients to derive a smoothed spectrum that is then perceptually weighted to give perceptually weighted linear prediction (PLP) coefficients [44].

As will be discussed in the next section, the acoustic modelling assumes that each acoustic vector is uncorrelated with its neighbors. This is a rather poor assumption since the physical constraints of the human vocal apparatus ensure that there is continuity between successive spectral estimates. However, appending the first- and second-order differentials to the basic static coefficients will greatly reduce the problem [3, 31]. In the HTK recognizer, these are approximated by fitting a linear regression over a window covering the two preceding and two following vectors. When this is done, the final acoustic vector has 39 components.

Although the above description is specific to one particular recognition system, it is typical of most modern LVR systems. An important point to emphasize is the degree to which the front end of modern recognizers has evolved to optimize the subsequent pattern matching. For example, in the above, the log compression, DCT transform, and delta coefficients are all introduced primarily to satisfy the assumptions made by the acoustic modelling component.

Acoustic Modelling

The purpose of the acoustic models is to provide a method of calculating the likelihood of any vector sequence Y given a

word w . In principle, the required probability distribution could be located by finding many examples of each w and collecting the statistics of the corresponding vector sequences. However, this is impractical for large-vocabulary systems and, instead, word sequences are decomposed into basic sounds called *phones*.

Each individual phone is represented by an HMM. An HMM has a number of states connected by arcs. HMM phone models typically have three *emitting* states and a simple left-right topology as illustrated by Fig. 3. The entry and exit states are provided to make it easy to join models together. The exit state of one phone model can be merged with the entry state of another to form a composite HMM. This allows phone models to be joined together to form words and words to be joined together to cover complete utterances.

An HMM is most easily understood as a generator of vector sequences. It is a finite-state machine that changes state once every time unit, and each time, t , that a state, j , is entered, an acoustic speech vector, y_t , is generated with probability density $b_j(y_t)$. Furthermore, the transition from state i to state j is also probabilistic and is governed by the discrete probability a_{ij} . Figure 3 shows an example of this process where the model moves through the state sequence $X = 1, 2, 2, 3, 4, 4, 5$ in order to generate the sequence y_1 to y_5 .

The joint probability of a vector sequence, Y , and state sequence X , given some model, M , is calculated simply as the product of the transition probabilities and the output probabilities. So for the state sequence X in Fig. 3:

$$P(Y, X|M) = a_{12}b_2(o_1)a_{22}b_2(o_2)a_{23}b_3(o_3) \dots \quad (2)$$

More formally, the joint probability of an acoustic vector sequence, Y , and some state sequence $X = x(1), x(2), x(3), \dots, x(T)$ is

$$P(Y, X|M) = a_{x(0)x(1)} \prod_{t=1}^T b_{x(t)}(y_t) a_{x(t)x(t+1)} \quad (3)$$

where $x(0)$ is constrained to be the model entry state and $x(T+1)$ is constrained to be the model exit state.

In practice, of course, only the observation sequence Y is known and the underlying state sequence X is hidden. This is why it is called a *hidden Markov model*. However, the required probability $P(Y|M)$ is easily found by summing Eq. 3 over all possible state sequences. The *Forward-Backward* algorithm is an efficient recursive method of performing this calculation. A crucial feature of this algorithm is that it also allows the probability of being in a specific model state at a specific time to be calculated. This leads to the *Baum-Welch* algorithm, which is a very simple and efficient procedure for finding maximum-likelihood estimates of both the a and b HMM parameter sets. Parameter estimation is beyond the scope of this article, but it is important to note that the existence of Baum-Welch has been a key factor in making HMMs the dominant technology in acoustic modelling.

Alternatively, $P(Y|M)$ can be approximated by finding the state sequence that maximizes Eq. 3. The *Viterbi* algorithm

is a simple algorithm for computing this efficiently. As will be discussed later, this algorithm is important in decoding, where determination of the most likely state sequence is the key to recognizing an unknown word sequence.

The brief outline of HMMs presented above is textbook material that has been well understood for many years. However, it is only recently that methods have been developed that allow HMM-based phone models to provide the acoustic discrimination necessary for large-vocabulary, speaker-independent speech recognition.

It is instructive to rewrite Eq. 3 in logarithmic form and separate out the a and b terms so that

$$\log P(Y, X|M) = \sum_{t=0}^T \log a_{x(t)x(t+1)} + \sum_{t=1}^T \log b_{x(t)}(y_t) \quad (4)$$


The transition probabilities $a_{x(t)x(t+1)}$ model the temporal structure of the data. Regarding each log probability in Eq. 4 as a *score*, each transition term can be viewed as the cost of moving from one state to another. This actually provides a very poor model for the duration of real speech, but this is not crucial since, in practice, the above expression is dominated by the output probabilities $b_{x(t)}(y_t)$. Each HMM state provides a prototype acoustic vector, and the log output probability function provides a distance metric to allow the actual acoustic vectors to be compared with the prototype.

The choice of output probability function is crucial since it must model all of the intrinsic spectral variability in real speech, both within and across speakers. Early HMM systems used discrete output probability functions in conjunction with a vector quantizer. Each incoming acoustic vector was replaced by the index of the closest vector in a precomputed codebook, and the output probability functions were just look-up tables containing the probabilities of each possible VQ index. Computationally, this approach is very efficient, but the quantization introduces noise that limits the precision that can be obtained. Hence, modern systems use parametric continuous-density output distributions that model the acoustic vectors directly [10, 56, 67]. The most common choice of distribution is the multivariate mixture Gaussian:

$$b_j(y_t) = \sum_{m=1}^M c_{jm} N(y_t; \mu_{jm}, \Sigma_{jm}) \quad (5)$$

where c_{jm} is the weight of mixture component m in state j and $N(y; \mu, \Sigma)$ denotes a multivariate Gaussian of mean μ and covariance Σ .

So far there has been an implicit assumption that only one HMM is required per phone, and since approximately 45 phones are needed for English, it may be thought that only 45 phone HMMs need be trained. In practice, however, contextual effects cause large variations in the way that different sounds are produced. Hence, to achieve good phonetic discrimination, different HMMs have to be trained for each different context. The simplest and most common approach is to use *triphones*, where every phone has a distinct HMM model for every unique pair of left and right neighbors. For example, suppose that the notation $x-y+z$ represents the



4. State Tying. In order to maximize the amount of data available to train each state while preserving discrimination ability, similar HMM states of the allophonic variants of each basic phone are tied together. The choice of which states to tie is made using a decision tree.

phone *y* occurring after an *x* and before a *z*. The phrase, "Beat it!" would be represented by the phone sequence **sil b iy t ih t sil**, and if triphone HMMs were used the sequence would be modelled as

sil sil-b+iy b-iy+t iy-t+ih t-ih+t ih-t+sil sil

Notice that the triphone contexts span word boundaries and the two instances of the phone *t* are represented by different HMMs because their contexts are different. This use of so-called *cross-word triphones* gives the best modelling accuracy, but it leads to complications in the decoder, as discussed later. Simpler systems result from the use of *word-internal triphones*, where the above example would become

sil b+iy b-iy+t iy-t ih+t ih-t sil

Here far fewer distinct models are needed, which simplifies both the parameter estimation problem and decoder design. However, the cost is an inability to model contextual effects at word boundaries, which, in fluent speech, are considerable.


The use of Gaussian mixture output distributions allows each state distribution to be modelled very accurately. However, when triphones are used they result in a system that has too many parameters to train. For example, a large-vocabulary cross-word triphone system will typically need around 60,000 triphones (With 45 phones, there are $45^3 = 91125$ possible triphones, but not all can occur due to the phonotactic constraints of the language). In practice, around 10 mixture components give good performance in LVR systems. Assuming that the covariances are all diagonal, then the HTK

recognizer with its 39-element acoustic vectors would require around 790 parameters per state. Hence, 60,000 3-state triphones would have a total of 142 million parameters.

The problem of too many parameters and too little training data is absolutely crucial in the design of a statistical speech recognizer. Early systems dealt with the problem by tying all Gaussian components together to form a pool that was then shared among all HMM states. In these so-called tied-mixture systems, only the mixture component weights were state specific, and these could be smoothed by interpolating with context-independent models [16, 48, 86]. Comparisons between discrete, tied-mixture, and continuous-density HMMs showed that tied-mixture HMMs were superior [47]. However, this followed mainly from the lack of good smoothing techniques for continuous-density systems. More recently, smoothing based on parameter tying has become popular [104]. In particular, state-tying [49, 106] and phone-based component tying [26] have been studied. Using these tying techniques with continuous-density HMMs has led to substantial improvements in modelling accuracy.

The HTK recognizer uses state tying. The idea is to tie together states that are acoustically indistinguishable. This allows all the data associated with each individual state to be pooled and thereby give more robust estimates for the parameters of the tied state. This is illustrated in Fig. 4. At the top of the figure, each triphone has its own private output distribution. After tying, several states share distributions.

In the HTK recognizer, the choice of which states to tie is made using *phonetic decision trees* [12, 57, 105]. This involves building a binary tree for each phone and state posi-



5. Decision Tree Clustering. Initially, all corresponding HMM states of all allophonic variants of each basic phone are tied to form a single pool. Phonetic questions are then used to partition the pool into subsets in a way that maximizes the likelihood of the training data. The leaf nodes of each tree determine the sets of state tyings for each of the allophonic variants.

tion. Each tree has a yes/no phonetic question such as “Is the left context a nasal?” at each node. Initially, all states for a given phone state position are placed at the root node of a tree. Depending on each answer, the pool of states is successively split and this continues until the states have trickled down to leaf nodes. All states in the same leaf node are then tied. For example, Fig. 5 illustrates the case of tying the center states of all triphones of the phone /aw/ (as in “out”). All of the states trickle down the tree, and, depending on the answer to the questions, they end up at one of the shaded terminal nodes. For example, in the illustrated case, the center state of **s-aw+n** would join the second leaf node from the right since its right context is a central consonant and a nasal, but its left context is not a central stop.

The questions at each node are chosen to maximize the likelihood of the training data given the final set of state tyings. In practice, phonetic decision trees give compact, good-quality state clusters that have sufficient associated data to robustly estimate mixture Gaussian output probability functions. Furthermore, they can be used to synthesize an HMM for any possible context, whether it appears in the training data or not, simply by descending the trees and using the state distributions associated with the terminating leaf nodes. Finally, phonetic decision trees can be used to include more than simple triphone contexts. For example, the HTK recognizer can use questions spanning ± 2 phones and can also take account of the presence of word boundaries.

Language Modelling

The purpose of the language model is to provide a mechanism for estimating the probability of some word, w_k , in an utterance given the preceding words, $W_1^{k-1} = w_1 \dots w_{k-1}$. A simple but effective way of doing this is to use N-grams, in which it is assumed that w_k depends only on the preceding $n - 1$ words, that is


$$P(w_k | W_1^{k-1}) = P(w_k | W_{k-n+1}^{k-1}) \quad (6)$$

N-grams simultaneously encode syntax, semantics, and pragmatics and they concentrate on local dependencies. This makes them very effective for languages like English where word order is important and the strongest contextual effects tend to come from near neighbors. Furthermore, N-gram probability distributions can be computed directly from text data and hence there is no requirement to have explicit linguistic rules such as a formal grammar of the language.

In principle, N-grams can be estimated from simple frequency counts and stored in a look-up table. For example, for the case of *trigrams* ($N = 3$),

$$\hat{P}(w_k | w_{k-1}, w_{k-2}) = \frac{t(w_{k-2}, w_{k-1}, w_k)}{b(w_{k-2}, w_{k-1})} \quad (7)$$

where $t(a, b, c)$ is the number of times the trigram a, b, c appears in the training data and $b(a, b)$ is the number of times the bigram a, b appears. The problem, of course, is that for a



6. Fragment of Decoder Network. In principle, the decoder searches through a network representing all possible word sequences. In practice, only paths corresponding to the most likely word sequences are constructed. Part (a) shows the directed network of words that the recognizer is considering initially. Part (b) shows the same network decomposed into triphones. Note that in order to take account of cross-word context, the first /a/ sound has to be replicated and the word *a* duplicated. Part (c) shows that tree-structuring the network can reduce the size of the network.

vocabulary of V words, there are V^3 potential trigrams. Even for a very modest vocabulary of 10,000 words, this is a very large number. Thus, many trigrams will not appear in the training data and many others will only appear once or twice, so that the estimate given by Eq. 7 will be very poor. In short, there is an acute data sparsity problem.

The solution to training data sparsity is to use a combination of *discounting* and *backing-off* [58, 77]. Discounting means that the trigram counts of the more frequently occurring trigrams are reduced and the resulting excess *probability mass* is redistributed among the less frequently occurring trigrams. Backing-off is applied when there are too few trigrams to form any estimate at all (e.g., just one or two occurrences in the training data). It involves replacing the trigram probability by a scaled bigram probability, that is

$$\hat{P}(w_k | w_{k-1}, w_{k-2}) = B(w_{k-1}, w_{k-2}) P(w_k | w_{k-1}) \quad (8)$$

where B is a back-off function included to ensure that $\hat{P}(w_k | w_{k-1}, w_{k-2})$ is properly normalized.

Although robust estimation of trigram probabilities requires considerable care, the problems are soluble and good performance can be obtained. N-grams do have obvious deficiencies resulting from their inability to exploit long-range constraints such as subject-verb agreement [55]. As a consequence, various alternatives have been studied such as tree-based models [11], trellis models [98], trigger models [63], history models [17], and variable N-grams [25]. However, in general, all of these attempts have yielded only small improvements at considerable computational cost. Thus to date, bigram and trigram language models dominate LVR systems.

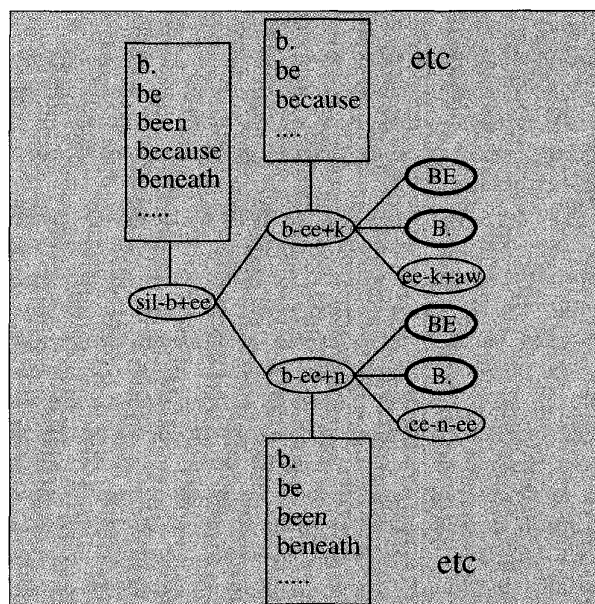
Decoding

The preceding sections have described the main components of a large-vocabulary system. In order to perform recognition using these components, the sequence of words \hat{W} that maximizes Eq. 1 must be found. This is a search problem and its solution is the domain of the decoder.

As with all search problems, there are two main approaches: *depth-first* and *breadth-first*. In depth-first designs, the most promising hypothesis is pursued until the end of the speech is reached. Examples of depth-first decoders are *stack-decoders* and *A*-decoders* [58, 59, 86, 87]. More recently, a refinement of stack decoding based on an *envelope search* has been proposed [40].

In breadth-first designs, all hypotheses are pursued in parallel. Breadth-first decoding exploits Bellman's optimality principle and is often referred to as *Viterbi decoding*. Since LVR systems are complex and pruning of the search space is essential, a process called *beam search* is typically used [41, 96]. The HTK decoder uses beam search and Viterbi decoding.

To understand the decoding problem, imagine that a branching tree network is constructed such that at the start there is a branch to every possible start word. All first words are then connected to all possible follow words and so on. This is illustrated in part (a) of Fig. 6. Clearly this tree will be very large, but if extended deep enough it would, in principle, represent all possible sequences. At first sight, this representation might seem very extravagant. In small-vocabulary systems, it is usually sufficient to put all words in parallel and place a loop around them. This allows all possible word sequences to be represented in a compact way since every vocabulary word appears only once. Unfortunately, however, such an arrangement does not allow a trigram language model to be used since the available history is limited to one word. Furthermore, a single loop back prevents cross-word triphones from being used. An explicit branching tree, however, allows both to be used in a straightforward manner [5].



7. Early Application of the Language Model. In a tree-structured network, the language model probability cannot be applied until the end of the word is reached. However, this delayed application severely limits the effectiveness of the language model for pruning. To solve this problem, each phone model carries a list of all possible words that it can belong to. This allows the probability of the most likely word to be used as an estimate for the probability of the actual word.

Next, let each word in this tree be replaced by the sequence of models representing its pronunciation. If there are multiple pronunciations then models can be joined in parallel within the word. Part (b) of Fig. 6 shows a fragment of the tree expanded into models. Finally, merge all identical phone models in identical contexts that have a common entry point as illustrated in part (c) of Fig. 6. Notice here that the use of cross-word triphones significantly limits the amount of model sharing possible.

The net result of the above is a branching tree of HMM-state nodes connected by state transitions and word-end nodes connected by word transitions. Any path from the start node to some point in the tree can be evaluated by adding all the log state transition probabilities, all the log state output probabilities, and the log language-model probabilities. Such a path can be represented by a movable *token* placed in the node at the end of the path [103]. The token has a *score*, which is the total log probability up to that point, and a *history*, which records the sequence of word-end nodes that the token has passed through. Any path can be extended by moving the token from its current node to an adjoining node and updating its score according to the current state-transition probability, state-output probability, and the language-model probability, if any.

The search problem can now be recast in the form of a token-passing algorithm. Initially, a single token is placed in the start node of the tree. As each acoustic vector is input, every token is copied into all connecting nodes and the scores updated. If more than one token lands in a node, only the best


scoring token needs to be retained since, by Bellmans' optimality principle, all other tokens must lie on inferior paths. When all of the acoustic vectors have been processed, the word-end nodes are scanned and the token with the highest score represents the most likely path and, hence, the most likely word sequence.

This basic token-passing algorithm is guaranteed to find the best possible path, but unfortunately, it would take too much time and space to compute. Hence, to make the algorithm tractable, pruning is employed. For every time frame, the best score in any token is noted and any token whose score lies more than a beam-width below this best score is destroyed. Since only the active tokens lying within the beam need to be kept in memory, only a fragment of the branching tree described above is ever needed at one time. As tokens move forward, new tree structure is created in front of them and the old structure behind them is destroyed. For this to work efficiently, it is crucial to prune tokens as soon as possible, and here tree-structuring causes a problem since a side-effect of merging phone models is that the identity of each new word entered is not known until its end node is reached. This is unfortunate since the language model provides a very powerful constraint, which needs to be applied as soon as is practicable in order to keep the number of active tokens as small as possible. The HTK decoder deals with this problem by associating a list of possible current words with every token (see Fig. 7). As the token moves toward the end of the word, this list will shrink until eventually it contains just a single word. Tokens then receive a language-model score equal to the most likely word in the current list. As this gets updated on every model transition, the tokens get pruned accordingly.

The dynamic network approach used in the HTK decoder results in a system that can arbitrarily exploit long-span language models and HMM phone models that depend on both the previous and succeeding acoustic context. Furthermore, it can do this in a single pass [81]. Most other Viterbi-based systems use a multiple-pass approach in which the first pass uses simple acoustic and language models and outputs a lattice of alternatives that are then rescored using more complex models in subsequent passes [6, 74, 90]. The problem with this is that an error in the first pass can never be recovered, hence large lattices must be used and any potential computational savings are lost. Multiple passes are useful for applying more complex language models, but for maximum

Table 1: HTK LVR Performance on ARPA CSR November 1994 Evaluation Data			
Vocab Size	LM N-gram	Adapted	% Word Error
20K	3	No	10.5†
65K	3	No	8.2
65K	4	No	7.9
65K	4	Yes	7.2‡

†H1-C1 test, ‡H1-P0 test



8. Word error rate plotted for each speaker using the HTK LVR (H1-P0) System on ARPA CSR November 1994 Evaluation Data (dashed horizontal line shows the average 7.2% WER).

accuracy, the most accurate acoustic models available need to be applied as soon as possible.

Current State of LVR

The major benchmarks for assessing the performance of LVR systems are the US Advanced Research Project Agency (ARPA) CSR Evaluations. The last full evaluation of dictation-style large-vocabulary recognition was the November 1994 evaluation [61, 85] in which the participating systems included AT&T Bell Laboratories [66, 69], BBN [78], Boston University [83], the CUED ABBOT group [46], the CUED HTK group [101, 100], IBM [8, 9], LIMSI-CNRS [38, 36], Philips [29], and SRI International [28].

The main focus of the evaluation was the so-called *hub test* H1 on which all participating sites evaluated their systems. This hub test was split into two main parts: H1-C1 in which the acoustic training data and a 20K-word trigram language model trained on 227 million words of news text were fixed; and H1-P0 in which any acoustic- or language-model training data could be used. In H1-C1 each utterance had to be recognized independently, whereas in H1-P0 each change of speaker was known so that unsupervised incremental adaptation could be used.

The HTK LV recognizer had the lowest error rate of the systems tested in the November 1994 evaluation, and it is therefore indicative of what can be achieved with current technology. Performance in terms of the percentage-word-error rate for a number of conditions (including H1-C1 and H1-P0) is shown in Table 1. As can be seen, the best performance achieved was 7.2 % (on average, seven words in every 100 were mistranscribed.) Although this figure is somewhat high, it is interesting to look at the error rates on a per-speaker basis as shown in Fig. 8, where the speakers have been ordered based on their performance. This figure suggests that, at least for part of the population, useable performance is achievable now. Conversely, it also shows that to cover the majority of the population, better robustness and more effective adaptation is needed (some training of the speakers would also help!).

Current Issues

The performance levels described in the previous section were all obtained for speech that was read in quiet recording conditions with a known microphone and a known task domain. Furthermore, the majority of systems tested were operating at many times slower than real time. Thus, although large-vocabulary continuous-speech recognition appears to be feasible, the following issues will need to be resolved before widespread deployment of LVR technology is possible.

Speaker Adaptation

Speaker independence is highly desirable since it allows a system to be used *straight out of the box*, and it allows systems to be built for which the speaker is not known in advance. However, in many applications, a speaker will either become a regular user or will input a reasonable quantity of speech at first use. In such cases it is natural to adapt the acoustic models to match that speaker. Adaptation can result in substantial reductions in error rate, particularly for atypical speakers, and its incorporation in future LVR systems will be essential.

Adaptation can be supervised or unsupervised, and it can be performed incrementally as the speaker is talking or offline at the end of the session. Each of these different styles may be most appropriate for some particular application. However, unsupervised incremental adaptation is the least intrusive and most generally useful to the user. Current research is therefore particularly concerned with techniques that can yield worthwhile performance gains with very little adaptation data, but which also asymptotically lead to speaker-dependent performance once a large amount of data has been acquired.

The acoustic models in an LVR system comprise a very large number of parameters. Hence, adaptation involving a very small amount of new data must depend on some form of general transformation rather than a direct re-estimation of the parameters themselves. In VQ or tied-mixture systems, this can be achieved by code-book mapping [95]. In CD systems, linear transformations can be applied to each Gaussian component. This can be a single global transform based on linear regression [45, 52] or canonical correlation [22]. As more data becomes available, Gaussians can be clustered and an individual transform determined for each cluster [27, 65]. Alternative approaches include clustering [32], regression-based prediction [23], and data augmentation, where the training data rather than the models is adapted [15]. Finally, when there is a reasonable amount of adaptation data, classical MAP estimation can be used to update all of the acoustic parameters [37, 64].

Environmental Robustness

Robustness to background noise and channel variability is clearly an essential requirement for the widespread use of

LVR technology. As explained above, LVR systems utilize acoustic feature vectors consisting of short-term spectral estimates to which some form of amplitude compression, such as a log operation, has been applied.

The primary effect of additive noise is to shift the spectral means and shrink the variances. However, it is important to note that the addition of noise changes the whole distribution and not just the means and variances [82]. Hence, a secondary effect of noise is a change to the modelling assumptions. Channel variability is convolutive noise that after a log operation becomes a simple, but possibly time-varying, offset.

There are a variety of approaches to dealing with noise [33, 42]. At the front-end, the noise can be removed from the speech [19], noise-robust features can be used [70, 73], the noise can be masked [71], or the features can be mapped [76]. The problem with these approaches is that, at best, they can only exploit knowledge of the global statistics of the speech in order to remove the noise. Given that the acoustic models within a recognizer encode very detailed information about the speech, this is an unnecessary handicap. Alternatively, an attempt can be made to make the pattern-matching process itself robust to noise [93]. However, in LVR systems, there are a very large number of overlapping acoustic classes, and maintaining maximal discrimination is essential.

Hence, for LVR systems, methods that adapt the recognizer to handle the corrupted speech signal directly are most attractive. These include code-book mapping for discrete and tied-mixture systems [2], cepstral mean compensation [72], state-based filtering [14], and parallel model combination (PMC) [34, 35, 97]. Note also that many of the techniques used for speaker adaptation are also capable of adapting to different noise environments.

In the large-vocabulary area, noise robustness remains a substantially unsolved problem. As the results in [85] show, without any form of compensation the performance of an LVR system will drop dramatically in noise. Compensation can limit this effect but cannot yet give immunity.

Task Independence

Whereas the acoustic models in an LVR system are relatively task independent, the language model is typically trained on a large corpus of task-specific material. This leads to systems that are inherently task dependent. For example, a language model trained on office correspondence will not work well on legal documents. Moving domain typically results in a lowering of accuracy due to the language-model mismatch and an increase in the incidence of out-of-vocabulary (OOV) words.

In the long term, task dependence will be reduced through a general increase in language-modelling capability. A recent trend in computational linguistics has been an increased interest in statistical techniques [21], and this should eventually lead to the incorporation of explicit grammatical knowledge into LVR systems. In the nearer term, solutions being studied include multiple-domain modelling and adaptation. In the former, individual LMs are trained on a range of

possible domains and then combined as a mixture model [50]. Adaptive systems typically combine a static LM trained off-line with a dynamic cache of N-grams that is continuously updated [62, 92].

The OOV problem has many facets [102], but increased robustness and the ability to simply add new words to an existing LM suggest that N-grams should be class-based rather than word-based. Using word classes allows compact, robust language models to be developed that can support very large vocabularies [20, 51, 60]. However, there are a number of problems to solve in using class-based N-grams. First, the method of choosing the classes is crucial since, ideally, words should only be grouped into a class if there is little or no resulting increase in perplexity. This implies that efficient, data-driven methods are needed. Second, since words can have multiple senses, it would be natural to allow the same word to appear in multiple classes. However, this adds considerable complexity to the decoding problem.

Spontaneous Speech

Much of the recent research effort in the LVR area has been directed at transcribing read speech. However, the ability to transcribe spontaneous casual speech is also an important requirement and this is now receiving increased attention. Current experimental evidence, particularly using the Switchboard Database, suggests that the error rate of current systems increases substantially when applied to spontaneous speech [107]. The reasons for this are still unclear but probably stem from a number of factors, including poor articulation, increased coarticulation, highly variable speaking rate, and various types of disfluency such as hesitations, false-starts, and corrections.

Some of the linguistic aspects of spontaneous speech have been studied in detail [84], and a number of specific issues such as identifying repairs in speech [43], modelling non-speech interjections [94], detecting hesitations [79], and modelling the timing patterns of disfluent speech [80] have been addressed. However, much of this work is at a very early stage. Practical implementations to date have been focussed on small- to medium-vocabulary dialogue systems (e.g., [99]) and little has been done in the large-vocabulary transcription area.

Real-Time Operation

In addition to obtaining an acceptable level of recognition accuracy, computationally efficient implementation is needed in order to exploit LVR technology. As explained in the section on decoding, recognition in LVR systems involves searching a very large space of hypotheses. In continuous-density systems, the computational cost of this search will be divided between building and searching the network structure and evaluating Gaussian densities.

The search space can be reduced by making approximations that allow alternative paths to be merged, for example, by approximating the language model [75] or by limiting the

number of new words through some form of look-ahead based on a fast-match preselection of possible followers [7, 39]. The evaluation of state output probabilities can be reduced by using vector-quantization to preselect those Gaussians that will give sufficiently high likelihoods [18]. Also, careful coding can make a substantial difference on modern RISC-based processor architectures where efficient use of the cache can make a substantial difference to throughput [1].

Conclusions

This article has reviewed the main components of a speaker-independent, continuous-speech LVR system and briefly described the state-of-the-art. While it is clear that much more needs to be done before robust, general-purpose LVR is ubiquitous, the technology is nevertheless on the threshold of usefulness for practical applications. Given a reasonably controlled environment and a well-defined task domain, the technology is useable now. By the end of 1996, off-the-shelf LVR systems that run in real-time on high-end, PC-class machines will start to appear. LVR-based services will also appear in the form of remote servers in public telecom systems, which will add telephone-based transcription capability to information and personal management services. Finally, LVR systems will facilitate multimedia information retrieval by allowing video soundtracks to be transcribed and then searched for key words and phrases.

Acknowledgments

Many people have contributed to the HTK LVR system, but the contributions of Phil Woodland, Julian Odell, and Valtcho Valtchev deserve particular mention. Also, the financial support of the UK Engineering and Physical Sciences Research Council (grant refs GR/J10204 and GR/K25380) is gratefully acknowledged.

Steve Young is a Professor of Information Engineering at the Cambridge University Engineering Department, Cambridge, UK.

References

1. F.A. Alleva. "Search Organisation for Large Vocabulary Continuous Speech Recognition." *Proc NATO Workshop*, 1990.
2. A. Anastasakos, F. Kubala, J. Makhoul, R. Schwartz. "Adaptation to New Microphones Using Tied-Mixture Normalisation." *Proc Human Language Technology Workshop*, pp. 325-329, Plainsboro NJ, Morgan Kaufman Publishers Inc, March, 1994.
3. T. Applebaum, B. Hanson. "Regression Features for Recognition of Speech in Noise." *Proc ICASSP*, S14.26, Toronto, 1991.
4. B.S. Atal. "Effectiveness of Linear Prediction Characteristics of the Speech Wave for Automatic Speaker Identification and Verification." *J Acoustical Soc Am*, Vol 55, No 6, pp. 1304-1312, 1974.
5. X. Aubert, H. Ney. "Large Vocabulary Continuous Speech Recognition Using Word Graphs." *Proc ICASSP*, Vol 1, pp. 49-52, Detroit, 1995.
6. S. Austin, R. Schwartz. "The Forward-Backward Search Algorithm for Continuous Speech Recognition." *Proc ICASSP*, S10.3, Toronto, 1991.

7. L. Bahl, P.S. Gopalakrishnan, D.S. Kanevsky, D. Nahamoo. "A Fast Admissible Method for Identifying a Short List of Candidate Words." *Computer Speech and Language*, Vol 6, No 3, pp. 215-224, 1992.
8. L.R. Bahl, S. Balakrishnan-Aiyer, J.R. Bellegarda, M. Franz, P.S. Gopalakrishnan, D. Nahamoo, M. Novak, M. Padmanabhan, M.A. Picheny, S. Roukos. "Performance of the IBM Large Vocabulary Continuous Speech Recognition System on the ARPA Wall Street Journal Task." *Proc ICASSP*, Vol 1, pp. 41-44, Detroit, 1995.
9. L.R. Bahl, S. Balakrishnan-Aiyer, M. Franz, P.S. Gopalakrishnan, R. Gopinath, M. Novak, M. Padmanabhan, S. Roukos. "The IBM Large Vocabulary Continuous Speech Recognition System for the ARPA NAB News Task." *Proc Spoken Language Technology Workshop*, pp. 121-126, Morgan Kaufmann Publishers Inc, Austin, Texas, Jan, 1995.
10. L.R. Bahl, P.F. Brown, P.V. de Souza, R.L. Mercer. "Speech Recognition with Continuous Parameters Hidden Markov Models." *Computer Speech and Language*, Vol 2, No 3/4, pp. 219-234, 1987.
11. L.R. Bahl, P.F. Brown, P.V. de Souza, R.L. Mercer. "A Tree-Based Statistical Language Model for Natural Language Speech Recognition." *IEEE Trans ASSP*, Vol 37, No 7, 1989.
12. L.R. Bahl, P.V. de Souza, P.S. Gopalakrishnan, D. Nahamoo, M.A. Picheny. "Context Dependent Modeling of Phones in Continuous Speech Using Decision Trees." *Proc DARPA Speech and Natural Language Processing Workshop*, pp. 264-270, Pacific Grove, Calif, Feb, 1991.
13. J.K. Baker. "The Dragon System - an Overview." *IEEE Trans ASSP*, Vol 23, No 1, pp. 24-29, 1975.
14. V.L. Beattie, S.J. Young. "Hidden Markov Model State-Based Cepstral Noise Compensation." *Proc ICSLP*, pp. 519-522, Banff, Canada, Oct, 1992.
15. J.R. Bellegarda, P.V. de Souza, D. Nahamoo, M. Padmanabhan, M.A. Picheny, L.R. Bahl. "Experiments using Data Augmentation for Speaker Adaptation." *Proc ICASSP*, Vol 1, pp. 692-695, Detroit, 1995.
16. J.R. Bellegarda, D. Nahamoo. "Tied Mixture Continuous Parameter Modeling for Speech Recognition." *IEEE Trans ASSP* Vol 38, No 12, pp.2033-2045, 1990.
17. E. Black, F. Jelinek, J. Lafferty, D.M. Magerman, R. Mercer, S. Roukos. "Towards History-based Grammars: Using Richer Models for Probabilistic Parsing." *Proc DARPA, Spoken Language Systems Workshop*, Feb, 1992.
18. E. Bocchieri. "A Study of the Beam Search Algorithm for Large Vocabulary Continuous Speech Recognition and Methods for Improved Efficiency." *Proc Eurospeech*, pp. 1521-1524, Berlin, 1993.
19. S. Boll. "Suppression of Acoustic Noise in Speech Using Spectral Subtraction." *IEEE Trans ASSP*, Vol 27, pp. 113-120, 1979.
20. P.F. Brown, V.J. Della Pietra, P.V. de Souza, J.C. Lai, R.L. Mercer. "Class-based n-gram Models of Natural Language." *Computational Linguistics*, Vol 18, No 4, pp. 467-479, 1992.
21. E. Charniak. *Statistical Language Learning*. Bradford MIT Press, 1993.
22. K. Choukri, G. Chollet. "Adaptation of Automatic Speech Recognizers to New Speakers Using Canonical Correlation Analysis Techniques." *Computer Speech Language*, Vol 1, No 2, pp. 95-107, 1986.
23. S. Cox. "Predictive Speaker Adaptation in Speech Recognition." *Computer Speech Language*, Vol 9, No 1, pp. 1-18, 1995.
24. S.B. Davis, P. Mermelstein. "Comparison of Parametric Representations for Monosyllabic Word Recognition in Continuously Spoken Sentences." *IEEE Trans ASSP*, Vol 28, No 4, pp. 357-366, 1980.
25. S. Deligne, F. Bimbot. "Language Modeling by Variable Length Sequences: Theoretical Formulation and Evaluation of Multigrams." *Proc ICASSP*, Vol 1, pp. 169-172, Detroit, 1995.
26. V. Digalakis, P. Monaco, H. Murveit. "Genones: Generalised Mixture Tying in Continuous Speech HMM-based Speech Recognizers." *IEEE Trans Speech and Audio Processing*, to appear, 1996.
27. V. Digalakis, D. Rtischev, L. Neumeyer. "Fast Speaker Adaptation Using Constrained Estimation of Gaussian Mixtures." *IEEE Trans on Speech and Audio Processing*, Vol 3, No 5, pp. 357-366, 1995.
28. V. Digilakis, M. Weintraub, Ananth Sankar, H. Franco. "Continuous Speech Dictation on ARPA's North American Business News Domain." *Proc Spoken Language Technology Workshop*, pp. 88-93, Morgan Kaufmann Publishers Inc, Austin, Texas, Jan, 1995.
29. C. Dugast, R. Kneser, X. Aubert, S. Ortmanns, K. Beulen, H. Ney. "Continuous Speech Recognition Tests and Results for the NAB'94 Corpus." *Proc Spoken Language Technology Workshop*, pp. 156-161, Morgan Kaufmann Publishers Inc, Austin, Texas, Jan, 1995.
30. S. Furui. "Cepstral Analysis Technique for Automatic Speaker Verification." *IEEE Trans ASSP*, Vol 29, No 2, pp. 254-272, 1981.
31. S. Furui. "Speaker-Independent Isolated Word Recognition Using Dynamic Features of Speech Spectrum." *IEEE Trans ASSP*, Vol 34, No 1, pp. 52-59, 1986.
32. S. Furui. "Unsupervised Speaker Adaptation Method Based on Hierarchical Clustering." *Proc ICASSP*, pp. 286-289, Glasgow, May, 1989.
33. S. Furui. "Toward Robust Speech Recognition under Adverse Conditions." *Proc ESCA Workshop on Speech Processing in Adverse Conditions*, pp. 31-42, Cannes-Mandelieu, November, 1992.
34. M.J.F. Gales, S.J. Young. "Cepstral Parameter Compensation for HMM Recognition in Noise." *Speech Communication*, Vol 12, No 3, pp. 231-239, 1993.
35. M.J.F. Gales, S.J. Young. "Robust Speech Recognition in Additive and Convolutional Noise using Parallel Model Combination." *Computer Speech and Language*, Vol 9, pp. 289-308, 1995.
36. J-L. Gauvain, L. Lamel, M. Adda-Decker. "Developments in Large Vocabulary Dictation: The LIMSI Nov94 NAB System." *Proc Spoken Language Technology Workshop*, pp. 131-138, Morgan Kaufmann Publishers Inc, Austin, Texas, Jan, 1995.
37. J-L. Gauvain, C-H Lee. "Maximum a Posteriori Estimation of Multivariate Gaussian Mixture Observations of Markov Chains." *IEEE Trans Speech and Audio Processing*, Vol 2, No 2, pp. 291-298, April, 1994.
38. J.L. Gauvain, L.F. Lamel, G. Adda, M. Adda-Decker. "The LIMSI Continuous Speech Dictation System." *Proc Human Language Technology Workshop*, pp. 319-324, Plainsboro NJ, Morgan Kaufman Publishers Inc, March, 1994.
39. L. Gillick, R. Roth. "A Rapid Match Algorithm for Continuous Speech Recognition." *Proc DARPA Speech and Natural Language Workshop*, Hidden Valley, Pennsylvania, pp. 170-172, June, 1990.
40. P.S. Gopalakrishnan, L.R. Bahl, R.L. Mercer. "A Tree Search Strategy for Large Vocabulary Continuous Speech Recognition." *Proc ICASSP*, Vol 1, pp. 572-575, Detroit, 1995.
41. R. Haeb-Umbach, H. Ney. "Improvements in Time-Synchronous Beam Search for 10000-Word Continuous Speech Recognition." *IEEE Trans Speech and Audio Processing*, Vol 2, pp. 353-356, 1994.
42. J-P. Haton. *Automatic Recognition of Speech in Adverse Conditions: a Review*. Reviewed for *IEEE Speech and Audio*, SAP 276, 1994.
43. P.A. Heeman, J. Allen. "Tagging Speech Repairs." *Proc Human Language Technology Workshop*, pp. 187-192, Plainsboro NJ, Morgan Kaufman Publishers Inc, March, 1994.
44. H. Hermansky. "Perceptual Linear Predictive (PLP) Analysis of Speech." *J Acoustical Soc America*, Vol 87, No 4, 1990.
45. A.J. Hewett, G. Holmes, S.J. Young. "Dynamic Speaker Adaptation in Speaker-Independent Word Recognition." *Proc IOA Autumn Conf*, Vol 8, Pt 7, pp. 275-282, 1986.
46. M. Hochberg, S. Renals, A. Robinson. "ABBOT: The CUED Hybrid Connectionist-HMM Large Vocabulary Recognition System." *Proc Spoken Language Technology Workshop*, pp. 170-178, Morgan Kaufmann Publishers Inc, Austin, Texas, Jan, 1995.
47. X.D. Huang, H.W. Hon, M.Y. Hwang, K.F. Lee. "A Comparative Study of Discrete, Semi-Continuous and Continuous Hidden Markov Models." *Computer Speech and Language*, Vol 7, No 4, pp. 359-368, 1993.
48. X.D. Huang, M.A. Jack. "Semi-continuous Hidden Markov Models for Speech Signals." *Computer Speech and Language*, Vol 3, No 3, pp. 239-252, 1989.
49. Hwang M-Y, X. Huang. "Shared Distribution Hidden Markov Models for Speech Recognition." *IEEE Trans Speech and Audio Processing*, Vol 1, No 4, pp. 414-420, 1993.

50. R. Iyer, M. Ostendorf, J.R. Rohlicek. "Language Modeling with Sentence Level Mixtures." *Proc Human Language Technology Workshop*, pp. 82-87, Plainsboro NJ, Morgan Kaufman Publishers Inc, March, 1994.
51. M. Jardino, G. Adda. "Automatic Word Classification using Simulated Annealing Language Modelling." *Proc ICASSP'93*, Vol 2, pp. 41-44, Minneapolis, 1993.
52. J. Janschul. "Speaker Adaptation by a Linear Transformation with Optimised Parameters." *Proc ICASSP*, pp. 1657-1660, Paris, 1982.
53. F. Jelinek. "A Fast Sequential Decoding Algorithm Using a Stack." *IBM J Research and Dev*, Vol 13, Nov, 1969.
54. F. Jelinek. "Continuous Speech Recognition by Statistical Methods." *Proc IEEE*, Vol 64, 4, pp. 532-556, 1976.
55. F. Jelinek. "Up From Trigrams: the Struggle for Improved Language Models." *Proc Eurospeech*, pp. 1037-1040, Genoa, 1991.
56. B-H Juang. "Maximum-Likelihood Estimation for Mixture Multivariate Stochastic Observations of Markov Chains." *AT&T Technical J*, Vol 64, No 6, pp. 1235-1249, 1985.
57. A. Kannan, M. Ostendorf, J.R. Rohlicek. "Maximum Likelihood Clustering of Gaussians for Speech Recognition." *IEEE Trans on Speech and Audio Processing*, Vol 2, No 3, pp. 453-455, 1994.
58. S.M. Katz. "Estimation of Probabilities from Sparse Data for the Language Model Component of a Speech Recogniser." *IEEE Trans ASSP*, Vol 35, No 3, pp. 400-401, 1987.
59. P. Kenny. "A* Admissible Heuristics for Rapid Lexical Access." *Proc ICASSP*, S10.1, Toronto, 1991.
60. R. Kneser, H. Ney. "Improved Clustering Techniques for Class-based Statistical Language Modelling." *Proc Eurospeech*, pp. 973-976, Berlin, 1993.
61. F. Kubala. "Design of the 1994 CSR Benchmark Tests." *Proc Spoken Language Technology Workshop*, pp. 41-46, Morgan Kaufmann Publishers Inc, Austin, Texas, Jan, 1995.
62. R. Kuhn, R. De Mori. "A Cache-Based Natural Language Model for Speech Recognition." *IEEE Trans Pattern Analysis and Machine Intelligence*, Vol 12, No 6, pp. 570-583, 1990.
63. R. Lau, R. Rosenfeld, S. Roukos. "Trigger-based Language Models: a Maximum Entropy Approach." *Proc ICASSP'93*, Vol 2, pp. 45-48, Minneapolis, 1993.
64. C-H Lee, C-H Lin, B-H Juang. "A Study on Speaker Adaptation of the Parameters of Continuous Density Hidden Markov Models." *IEEE Trans ASSP*, Vol 39, No 4, 1991.
65. C.J. Leggetter, P.C. Woodland. "Maximum Likelihood Linear Regression for Speaker Adaptation of Continuous Density Hidden Markov Models." *Computer Speech and Language*, Vol 9, No 2, pp. 171-185, 1995.
66. S.E. Levinson, A. Ljolje, L.G. Miller. "Large Vocabulary Speech Recognition Using a Hidden Markov Model for Acoustic/Phonetic Classification." *Proc ICASSP*, S11.5, pp. 505-508, 1988.
67. L.A. Liporace. "Maximum-Likelihood Estimation for Multivariate Observations of Markov Sources." *IEEE Trans Information Th*, Vol IT-28, No 5, pp. 729-734, 1982.
68. A. Ljolje. "The Importance of Cepstral Parameter Correlations in Speech Recognition." *Computer Speech and Language*, Vol 8, No 3, pp. 223-232, 1994.
69. A. Ljolje, M. Riley, D. Hindle, F. Pereira. "The AT&T 60,000 Word Speech-To-Text System." *Proc Spoken Language Technology Workshop*, pp. 162-165, Morgan Kaufmann Publishers Inc, Austin, Texas, Jan, 1995.
70. D. Mansour, B-H Juang. "The Short-Time Modified Coherence Representation and Noisy Speech Recognition." *IEEE Trans ASSP*, Vol 37, No 6, June, 1989.
71. B.A. Mellor, A.P. Varga. "Noise Masking in the MFCC Domain for the Recognition of Speech in Background Noise." *Proc Inst Acoustics Autumn Conf on Speech and Hearing*, Vol 14, Part 6, pp. 361-368, 1992.
72. P.J. Moreno, B. Raj, E. Gouvea, R.M. Stern. "Multivariate-Gaussian-Based Cepstral Normalisation for Robust Speech Recognition." *Proc ICASSP*, Vol 1, pp. 137-140, Detroit, 1995.
73. N. Morgan, H. Hermansky. "RASTA Extensions: Robustness to Additive and Convolutional Noise." *Proc ESCA Workshop on Speech Processing in Adverse Conditions*, pp. 115-118, Cannes-Mandelieu, November, 1992.
74. H. Murveit, J. Butzberger, V. Digilakis, M. Weintraub. "Large Vocabulary Dictation Using SRI's Decipher Speech Recognition System: Progressive Search Techniques." *Proc ICASSP*, pp. 319-322, Minneapolis, 1993.
75. H. Murveit, P. Monaco, V. Digilakis, J. Butzberger. "Techniques to Achieve an Accurate Real-Time Large Vocabulary Speech Recognition System." *Proc Human Language Technology Workshop*, pp. 393-398, Plainsboro NJ, Morgan Kaufman Publishers Inc, March, 1994.
76. L. Neumeyer, M. Weintraub. "Robust Speech Recognition in Noise Using Adaptation and Mapping Techniques." *Proc Spoken Language Technology Workshop*, pp. 100-103, Morgan Kaufmann Publishers Inc, Austin, Texas, Jan, 1995.
77. H. Ney, U. Essen, R. Kneser. "On Structuring Probabilistic Dependences in Stochastic Language Modelling." *Computer Speech and Language*, Vol 8, No 1, pp. 1-38, 1994.
78. L. Nguyen, Tasos Anastasakos, F. Kubala, C. LaPre, J. Makhoul, R. Schwartz, N. Yuan, G. Zavaliagkos, Y. Zhao. "The 1994 BBN BYBLOS Speech Recognition System." *Proc Spoken Language Technology Workshop*, pp. 77-81, Morgan Kaufmann Publishers Inc, Austin, Texas, Jan, 1995.
79. D. O'Shaughnessy. "Automatic Recognition of Hesitations in Spontaneous Speech." *Proc ICASSP*, Vol 1, pp. 593-596, San Francisco, 1992.
80. D. O'Shaughnessy. "Timing Patterns in Fluent and Disfluent Spontaneous Speech." *Proc ICASSP*, Vol 1, pp. 593-596, Detroit, 1995.
81. J.J. Odell, V. Valtchev, P.C. Woodland, S.J. Young. "A One-Pass Decoder Design for Large Vocabulary Recognition." *Proc Human Language Technology Workshop*, pp. 405-410, Plainsboro NJ, Morgan Kaufman Publishers Inc, March, 1994.
82. J.P. Openshaw, J.S. Mason. "On the Limitations of Cepstral Features in Noise." *Proc ICASSP*, pp. 49-52, Adelaide, Australia, April, 1994.
83. M. Ostendorf, F. Richardson, R. Iyer, A. Kannan, O. Ronen, R. Bates. "The 1994 BU NAB News Benchmark System." *Proc Spoken Language Technology Workshop*, pp. 139-142, Morgan Kaufmann Publishers Inc, Austin, Texas, Jan, 1995.
84. S. Oviatt. "Predicting Spoken Disfluencies during Human-Computer Interaction." *Computer Speech Language*, Vol 9, No 1, pp. 19-36, 1995.
85. D.S. Pallett, J.G. Fiscus, W.M. Fisher, J.S. Garofolo, B.S. Lund, A. Martin, Przybocki. "The 1994 Benchmark Tests for the ARPA Spoken Language Program." *Proc Spoken Language Technology Workshop*, pp. 5-38, Morgan Kaufmann Publishers Inc, Austin, Texas, Jan, 1995.
86. D.B. Paul. "The Lincoln Tied Mixture HMM Continuous Speech Recogniser." *Proc DARPA Speech and Natural Language Workshop*, pp. 332-336, Hidden Valley, Pennsylvania, June, 1990.
87. D.B. Paul. "Algorithms for an Optimal A* Search and Linearizing the Search in the Stack Decoder." *Proc ICASSP*, pp. 693-696, Toronto, 1991.
88. L.R. Rabiner, B-H Juang, S.E. Levinson, N.M. Sondhi. "Recognition of Isolated Digits Using HMMs with Continuous Mixture Densities." *AT&T Technical J*, Vol 64, No 6, pp. 1211-1233, 1985.
89. L.R. Rabiner, R.W. Schafer. *Digital Processing of Speech Signals*. Prentice-Hall, Englewood Cliffs, NJ, 1978.
90. F. Richardson, M. Ostendorf, J.R. Rohlicek. "Lattice-Based Search Strategies for Large Vocabulary Recognition." *Proc ICASSP*, Vol 1, pp. 576-579, Detroit, 1995.
91. A.J. Robinson, J. Holdsworth, R. Patterson, F. Fallside. "A Comparison of Preprocessors for the Cambridge Recurrent Error Propagation Network Speech Recognition System." *Proc International Conference on Spoken Language Processing*, Kobe, Japan, November, 1990.
92. R. Rosenfeld. "A Hybrid Approach to Adaptive Statistical Modeling." *Proc Human Language Technology Workshop*, pp. 76-81, Plainsboro NJ, Morgan Kaufman Publishers Inc, March, 1994.
93. A. Sankar, C-H Lee. "Robust Speech Recognition Based on Stochastic Matching." *Proc ICASSP*, Vol 1, pp. 121-124, Detroit, 1995.

94. T. Schultz, I. Rogina. "Acoustic and Language Modeling of Human and Nonhuman Noises for Human-Human Spontaneous Speech Recognition." *Proc ICASSP*, Vol 1, pp. 293-296, Detroit, 1995.
95. K. Shikano, K. Lee, R. Reddy. "Speaker adaptation through vector quantisation." *Proc ICASSP*, pp. 2643-2646, 1986.
96. V. Steinbiss. "Sentence-Hypotheses Generation in a Continuous-Speech Recognition System." *Proc European Conf on Speech Communication and Technology*, Vol 2, pp. 51-54, Paris, 1989.
97. A.P. Varga, R.K. Moore. "Hidden Markov Model Decomposition of Speech and Noise." *Proc ICASSP*, Albuquerque, 1990.
98. N.P. Waegner, S.J. Young. "A Trellis-based Language Model for Speech Recognition." *Proc ICSLP*, pp. 245-248, Banff, Canada, Oct, 1992.
99. W. Ward. "Understanding Spontaneous Speech: the Phoenix System." *Proc ICASSP 91*, pp. 365-368, Toronto, Canada, May, 1991.
100. P.C. Woodland, C.J. Leggetter, J. Odell, V. Valtchev, S.J. Young. "The 1994 HTK Large Vocabulary Speech Recognition System." *Proc ICASSP*, Vol 1, pp. 73-76, Detroit, 1995
101. P.C. Woodland, S.J. Young. "The HTK Continuous Speech Recogniser." *Proc Eurospeech '93*, pp. 2207-2219, Berlin, Sept, 1993.
102. S.R. Young. "Learning New Words from Spontaneous Speech: A Project Summary." *CMU Tech Report*, CMU-CS-93-223, July, 1993.
103. S.J. Young, N.H. Russell, J.H.S. Thornton. "Token Passing: A Simple Conceptual Model for Connected Speech Recognition Systems." *Technical Report CUED/F-INFENG/TR38*, Cambridge University Engineering Dept, 1989.
104. S.J. Young. "The General Use of Tying in Phoneme-Based HMM Speech Recognisers." *Proc ICASSP*, Vol 1, pp. 569-572, San Francisco, March, 1992.
105. S.J. Young, J.J. Odell, P.C. Woodland. "Tree-Based State Tying for High Accuracy Acoustic Modelling." *Proc Human Language Technology Workshop*, pp. 307-312, Plainsboro NJ, Morgan Kaufman Publishers Inc, March, 1994.
106. S.J. Young, P.C. Woodland. "State Clustering in HMM-based Continuous Speech Recognition." *Computer Speech and Language*, Vol 8, No 4, pp. 369-384, 1994.
107. S.J. Young, P.C. Woodland, W.J. Byrne. "Spontaneous Speech Recognition for the Credit Card Corpus using the HTK Toolkit." *IEEE Trans Audio and Speech Processing*, Vol 2, No 4, pp. 615-621, 1994.

DSP and **real-time** developments **on time** with the leading edge *Virtuoso*TM tools !

DSP

21020
21060
5600x
9600x
320C30
320C31
320C4x
320C5x
320C80
DSP cores

Other

80x86
3052
68HC11
68HC16
680x0
T4xx
T8xx
80960
ARM
PowerPC
SPARC
Alpha
...

Virtuoso's multi-level support

interrupt support
multi-tasking nanokernel
preemptive microkernel
transparent parallel processing
heterogeneous targets

No compromise on performance

ultrafast (ctxt switch < 1 µs)
small size (0.2 - 6.5 K Instr.)

Development support

system generation
(network) loader
debugger and tracing monitor
true portability and scalability
studio & PC graphics
cross development on PC
MS-Windows, Solaris integration
internet sockets
DSP & multi-media support
drivers
source code
12 months maintenance and
support included

Ease of use

Start on day one !

Virtuoso's multi-tool approach

✓ **Real-time O.S.'s :**

Virtuoso Micro
Virtuoso Classico
Virtuoso Nano

RTOS implementations :

SP : Single Processor
MP : Multi-Processor
VSP : Virtual Single Processor

✓ **Dataflow application generator :**

Virtuoso Synchro

✓ **DSP Libraries :**

Virtuoso Modulo
0, I, II, III, IV, V, VI

✓ **Services**

Customization and training
Consultancy for HW and SW
system design

Applications

signal processing
image processing
high speed control
telecommunications
data-acquisition
multi-media
military
space missions
embedded systems
distributed control
process control
...



Eonic Systems, Inc. 12210 Plum Orchard Drive, Silver Spring, MD 20904, USA, Tel. (301) 572 5000, Fax (301) 572 5005, e-mail: eonic@blix.com
Eonic Systems NV/SA, Lindesstraat 9, B-3210 Linden, Belgium, Tel. (+32) 16 62 15 85, Fax (+32) 16 62 15 84, e-mail: virtuoso@blix.com

Go VirtuosoTM !

Reader Service Number 23